

Statistics and Data Analysis Assignment

Domen Mohorčič, Larsen Cundrič

09/04/2021

Roller Coaster Dataset Analysis

State hypothesis here...

Basic analysis and plots

```
roller_coasters_raw <- readr::read_csv('datasets/roller_coasters.csv')
```

```
## Parsed with column specification:
## cols(
##   Name = col_character(),
##   Park = col_character(),
##   City = col_character(),
##   State = col_character(),
##   Country = col_character(),
##   Type = col_character(),
##   Construction = col_character(),
##   Height = col_double(),
##   Speed = col_double(),
##   Length = col_double(),
##   Inversions = col_character(),
##   Numinversions = col_double(),
##   Duration = col_double(),
##   GForce = col_double(),
##   Opened = col_double(),
##   Region = col_character()
## )
```

```
# GForce to many missing values..
```

```
#
```


```
summary(roller_coasters_raw)
```

```
##      Name      Park      City      State
## Length:408   Length:408   Length:408   Length:408
## Class :character Class :character Class :character Class :character
## Mode  :character Mode  :character Mode  :character Mode  :character
##
##
```

```
##
##
## Country           Type           Construction           Height
## Length:408        Length:408        Length:408          Min.   : 2.438
## Class :character   Class :character   Class :character    1st Qu.: 8.651
## Mode  :character   Mode  :character   Mode  :character    Median : 18.288
##                                     Mean   : 23.125
##                                     3rd Qu.: 33.167
##                                     Max.   :128.016
##                                     NA's   :82
## Speed            Length           Inversions           NumInversions
## Min.   : 9.72     Min.   : 12.19        Length:408          Min.   : 0.0000
## 1st Qu.: 45.00     1st Qu.: 291.00       Class :character    1st Qu.: 0.0000
## Median : 68.85     Median : 415.75       Mode  :character    Median : 0.0000
## Mean   : 69.36     Mean   : 597.04              Mean : 0.7843
## 3rd Qu.: 88.95     3rd Qu.: 833.12              3rd Qu.: 0.0000
## Max.   :194.40     Max.   :2243.02              Max.   :10.0000
## NA's   :138       NA's   :90
## Duration          GForce           Opened              Region
## Min.   : 0.3       Min.   :2.100         Min.   :1924        Length:408
## 1st Qu.: 75.0      1st Qu.:3.175         1st Qu.:1991        Class :character
## Median :108.0      Median :4.500         Median :1999        Mode  :character
## Mean   :112.5      Mean   :4.115         Mean   :1995
## 3rd Qu.:140.8      3rd Qu.:5.000         3rd Qu.:2004
## Max.   :300.0      Max.   :6.200         Max.   :2014
## NA's   :216       NA's   :348         NA's   :28
```


```
ggplot(data = roller_coasters_raw) +
  geom_boxplot(mapping = aes(x = Country, y = Speed))
```

```
## Warning: Removed 138 rows containing non-finite values (stat_boxplot).
```




```
ggplot(data = roller_coasters_raw) +  
  geom_boxplot(mapping = aes(x = City, y = Speed))
```

```
## Warning: Removed 138 rows containing non-finite values (stat_boxplot).
```




```
ggplot(data = roller_coasters_raw) +  
  geom_boxplot(mapping = aes(x = Park, y = Speed))
```

```
## Warning: Removed 138 rows containing non-finite values (stat_boxplot).
```




```
ggplot(data = roller_coasters_raw) +  
  geom_boxplot(mapping = aes(x = Construction, y = Speed))
```

```
## Warning: Removed 138 rows containing non-finite values (stat_boxplot).
```




```
# Type is same as Construction?  
ggplot(data = roller_coasters_raw) +  
  geom_boxplot(mapping = aes(x = Type, y = Speed))
```

```
## Warning: Removed 138 rows containing non-finite values (stat_boxplot).
```




```
roller_coasters_raw %>%  
  ggplot() +  
  geom_point(aes(x = Height, y = Speed, color = Construction))
```

```
## Warning: Removed 150 rows containing missing values (geom_point).
```




```
roller_coasters_raw %>%  
  ggplot() +  
  geom_point(aes(x = Length, y = Speed, color = Construction))
```

```
## Warning: Removed 148 rows containing missing values (geom_point).
```


```
roller_coasters_raw %>%  
  ggplot() +  
  geom_boxplot(aes(x = as.factor(Numinversions), y = Speed, color = Construction))
```

```
## Warning: Removed 138 rows containing non-finite values (stat_boxplot).
```



```
# dost lame ... lahko spustimo
roller_coasters_raw %>%
  filter(!is.na(GForce)) %>%
  ggplot() +
    geom_point(aes(x = GForce, y = Speed))
```

```
## Warning: Removed 2 rows containing missing values (geom_point).
```



Inference?

Correlation Analysis

```
# precej zanimivi so Height, Length, NumInversions
(cor.test(roller_coasters_raw$Height, roller_coasters_raw$Speed))
```

```
##
## Pearson's product-moment correlation
##
## data: roller_coasters_raw$Height and roller_coasters_raw$Speed
## t = 38.222, df = 256, p-value < 2.2e-16
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
##  0.9019179 0.9388051
## sample estimates:
##      cor
## 0.9224392
```

```
(cor.test(roller_coasters_raw$Length, roller_coasters_raw$Speed))
```

```
##
```

```

## Pearson's product-moment correlation
##
## data: roller_coasters_raw$Length and roller_coasters_raw$Speed
## t = 15.582, df = 258, p-value < 2.2e-16
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## 0.6278199 0.7540719
## sample estimates:
## cor
## 0.6962931

(cor.test(roller_coasters_raw$Numinversions, roller_coasters_raw$Speed))

##
## Pearson's product-moment correlation
##
## data: roller_coasters_raw$Numinversions and roller_coasters_raw$Speed
## t = 5.5742, df = 268, p-value = 6.061e-08
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## 0.2110692 0.4253337
## sample estimates:
## cor
## 0.3223236

(cor.test(roller_coasters_raw$Duration, roller_coasters_raw$Speed))

##
## Pearson's product-moment correlation
##
## data: roller_coasters_raw$Duration and roller_coasters_raw$Speed
## t = 3.9954, df = 162, p-value = 9.781e-05
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## 0.1532868 0.4328823
## sample estimates:
## cor
## 0.2995011

(cor.test(roller_coasters_raw$GForce, roller_coasters_raw$Speed))

##
## Pearson's product-moment correlation
##
## data: roller_coasters_raw$GForce and roller_coasters_raw$Speed
## t = 3.3676, df = 56, p-value = 0.001377
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## 0.1701111 0.6045861
## sample estimates:
## cor
## 0.4103754

```

```
(cor.test(roller_coasters_raw$Opened, roller_coasters_raw$Speed)) # shit
```

```
##  
## Pearson's product-moment correlation  
##  
## data: roller_coasters_raw$Opened and roller_coasters_raw$Speed  
## t = 0.26238, df = 260, p-value = 0.7932  
## alternative hypothesis: true correlation is not equal to 0  
## 95 percent confidence interval:  
## -0.1051251 0.1371870  
## sample estimates:  
## cor  
## 0.01626982
```

Regression

```
roller_coasters <- roller_coasters_raw %>%  
  select(Construction, Length, Height, Speed) %>%  
  filter(!is.na(Speed) & !is.na(Height) & !is.na(Length)) %>%  
  mutate("Steel" = as.numeric(Construction == 'Steel')) %>%  
  select(-Construction)  
roller_coasters
```

```
## # A tibble: 252 x 4  
##   Length Height Speed Steel  
##   <dbl> <dbl> <dbl> <dbl>  
## 1 853. 128. 194. 1  
## 2 376. 126. 162. 1  
## 3 2010. 94.5 151. 1  
## 4 1619. 74.7 138. 1  
## 5 1620 73 127 1  
## 6 1372. 71.6 138. 1  
## 7 1610. 70.1 130. 1  
## 8 1700 67 120 1  
## 9 2143. 66.4 126. 0  
## 10 396. 66.4 113. 1  
## # ... with 242 more rows
```

```
## 75% of the sample size  
smp_size <- floor(0.75 * nrow(roller_coasters))  
  
## set the seed to make your partition reproducible  
set.seed(123)  
train_ind <- sample(seq_len(nrow(roller_coasters)), size = smp_size)  
  
(train <- roller_coasters[train_ind, ])
```

```
## # A tibble: 189 x 4  
##   Length Height Speed Steel  
##   <dbl> <dbl> <dbl> <dbl>
```

```
## 1 412. 16.2 50.2 1
## 2 207 8.5 34.9 1
## 3 538. 13.4 50.2 1
## 4 375. 61 105. 1
## 5 427. 10.7 32.4 0
## 6 774. 14.6 68.0 1
## 7 950 36 90 1
## 8 717. 23.8 82.6 1
## 9 309. 39.9 81 1
## 10 264 6 45 1
## # ... with 179 more rows
```

```
(test <- roller_coasters[-train_ind, ])
```

```
## # A tibble: 63 x 4
##   Length Height Speed Steel
##   <dbl> <dbl> <dbl> <dbl>
## 1 376. 126. 162 1
## 2 2010. 94.5 151. 1
## 3 671. 62.5 133. 1
## 4 347. 61 105. 1
## 5 367. 58.2 105. 1
## 6 1167. 57.3 113. 1
## 7 1654. 49.1 105. 0
## 8 1332. 47.5 105. 1
## 9 150 46 105 1
## 10 192. 45.7 105. 1
## # ... with 53 more rows
```

```
lin_model <- lm(Speed ~ ., data = train)
(summary(lin_model))
```

```
##
## Call:
## lm(formula = Speed ~ ., data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -21.388  -6.020  -0.644   4.466  35.365
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 33.376471   2.513555  13.279 < 2e-16 ***
## Length      0.013174   0.002178   6.047 7.97e-09 ***
## Height      1.248969   0.047545  26.269 < 2e-16 ***
## Steel       -5.752860   2.109654  -2.727 0.00701 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.749 on 185 degrees of freedom
## Multiple R-squared:  0.9103, Adjusted R-squared:  0.9088
## F-statistic: 625.7 on 3 and 185 DF, p-value: < 2.2e-16
```

```
(coef(lin_model))
```

```
## (Intercept)      Length      Height      Steel  
## 33.37647051  0.01317372  1.24896948 -5.75286049
```

Second Dataset Analysis → TODO

State hypothesis here...